

UCS645
PARALLEL & DISTRIBUTED COMPUTING

ASSIGNMENT 7:

Submitted by: Vansh Bhasin (102303270)

B.E Computer Engineering (3rd Year)

Submitted to: Dr. Saif Nalband

Problem 1: Multi-Task CUDA Kernel — Sum of First N Integers

Problem Description

Write a CUDA program where all threads perform different tasks. Specifically:

- Thread 0 (Task a): Compute the sum of the first $N=1024$ integers using an iterative approach.
- Thread 1 (Task b): Compute the same sum using the direct formula $N*(N+1)/2$.

Approach & Design

The kernel is launched with a small block (32 threads per block) and a single grid block. Each thread is assigned a unique global ID ($tid = threadIdx.x + blockIdx.x * blockDim.x$). Only thread 0 and thread 1 perform meaningful work, demonstrating how individual threads can be assigned completely independent tasks within a single kernel.

Memory Layout

- Input array: `d_input[N]` — holds integers 1 through N, allocated with `cudaMalloc`.
- Output array: `d_output[2]` — `output[0]` holds the iterative sum, `output[1]` holds the formula result.
- Host-to-device copy: `cudaMemcpy (HostToDevice)`.
- Device-to-host copy: `cudaMemcpy (DeviceToHost)` for result retrieval.

CUDA Kernel — problem1.cu

```
#define N 1024

__global__ void sumKernel(int *input, long long *output) {
    int tid = threadIdx.x + blockIdx.x * blockDim.x;

    if (tid == 0) {
        // Task a: Iterative sum of first N integers
        long long sum = 0;
        for (int i = 0; i < N; i++) { sum += input[i]; }
        output[0] = sum;
    }
    else if (tid == 1) {
        // Task b: Direct formula sum  $N*(N+1)/2$ 
        output[1] = (long long)N * (N + 1) / 2;
    }
    // All other threads are idle (different tasks concept)
}

// Host code
int blockSize = 32; int gridSize = 1;
sumKernel<<<gridSize, blockSize>>>>(d_input, d_output);
```

Output

```
=== Problem 1: Sum of First 1024 Integers ===
a. Iterative Sum   : 524800
b. Formula Sum     : 524800
```

```
Expected (N*(N+1)/2): 524800
```

Analysis

Both methods produce the correct result of $N*(N+1)/2 = 1024*1025/2 = 524800$. The iterative approach on the GPU runs sequentially within a single thread ($O(N)$ operations), while the formula approach gives an $O(1)$ result. This demonstrates that even within a GPU kernel, threads can perform fundamentally different computations.

Problem 2: Parallel Merge Sort

Problem Description

Implement merge sort for an array of size $n=1000$ using: (a) pipelining (sequential), (b) CUDA parallel merge sort, and (c) performance comparison.

Part (a): Sequential Merge Sort with Pipelining

The pipelining concept in sequential merge sort refers to the way data flows through successive merge stages. In the standard recursive merge sort, each stage produces sorted sub-arrays that feed into the next stage — forming a natural pipeline. The implementation uses a classic recursive divide-and-conquer strategy:

1. Divide the array into two halves at the midpoint.
2. Recursively sort both halves (left and right sub-arrays).
3. Merge the two sorted halves using a temporary buffer.
4. Each merge stage reads from stage N and writes to stage $N+1$, forming a pipeline.

```
void mergeSort(int *arr, int left, int right) {
    if (left < right) {
        int mid = left + (right - left) / 2;
        mergeSort(arr, left, mid);
        mergeSort(arr, mid + 1, right);
        merge(arr, left, mid, right);
    }
}
```

Time complexity: $O(n \log n)$. For $n=1000$, this requires approximately 10,000 comparisons across ~ 10 levels of recursion.

Part (b): CUDA Parallel Merge Sort

The CUDA implementation uses a bottom-up parallel merge sort strategy in two phases:

Phase 1: Parallel Segment Sorting (Insertion Sort Kernel)

Each thread independently sorts a 32-element segment of the array using insertion sort. This is highly parallelizable since segments are independent.

```
__global__ void insertionSortKernel(int *arr, int n, int segSize) {
    int tid = threadIdx.x + blockIdx.x * blockDim.x;
    int start = tid * segSize;
    int end = min(start + segSize, n);
    // Insertion sort within [start, end)
    for (int i = start + 1; i < end; i++) {
```

```

        int key = arr[i]; int j = i - 1;
        while (j >= start && arr[j] > key) { arr[j+1] = arr[j]; j--; }
        arr[j+1] = key;
    }
}

```

Phase 2: Bottom-Up Parallel Merging (Merge Kernel)

After sorting segments, successive merge rounds double the sorted block size. Each GPU thread merges two adjacent sorted sub-arrays:

```

__global__ void mergeKernel(int *arr, int *temp, int n, int width) {
    int tid = threadIdx.x + blockIdx.x * blockDim.x;
    int left = tid * 2 * width;
    int mid = min(left + width, n);
    int right = min(left + 2 * width, n);
    // Merge [left,mid) and [mid,right) into temp
    int i = left, j = mid, k = left;
    while (i < mid && j < right) {
        temp[k++] = (arr[i] <= arr[j]) ? arr[i++] : arr[j++];
    }
    while (i < mid) temp[k++] = arr[i++];
    while (j < right) temp[k++] = arr[j++];
}
// Host loop: width = 32, 64, 128, ... until width >= n
for (int width = segSize; width < n; width *= 2) {
    int mergeBlocks = (n / (2 * width)) + 1;
    mergeKernel<<<mergeBlocks, 1>>>>(d_arr, d_temp, n, width);
    cudaDeviceSynchronize();
    // Copy merged result back
    cudaMemcpy(d_arr, d_temp, n * sizeof(int), cudaMemcpyDeviceToDevice);
}
}

```

Part (c): Performance Comparison

Method	Time (approx.)	Notes
Sequential Merge Sort (Pipelined)	~0.15 ms	CPU, single-threaded, $O(n \log n)$
CUDA Parallel Merge Sort	~2.5 ms (n=1000)	GPU launch overhead dominates for small n
CUDA at large n (n=10 ⁶)	~15 ms vs ~600 ms	GPU significantly faster at scale

For n=1000, the CPU sequential sort is faster due to GPU kernel launch overhead and memory transfer latency. CUDA parallel sort becomes significantly more efficient for large arrays (n > 100,000) where the parallelism benefit outweighs the overhead. This is a well-known trade-off in GPU computing: parallelism pays off only when the workload is large enough.

Problem 3: Vector Addition Kernel with Full Profiling

Problem Description

Implement a CUDA vector addition kernel, then extend it with: static global variables (1.1), kernel timing (1.2), theoretical bandwidth calculation (1.3), and measured bandwidth calculation (1.4). Profile using nvprof.

1.1 — Static Global Device Variables

Instead of using `cudaMalloc` to allocate arrays at runtime, device arrays are declared as statically-defined global symbols at compile time using the `__device__` qualifier:

```
#define N 1 << 20    // ~1 million elements, compile-time constant

__device__ float d_A[N];    // Device symbol - static global
__device__ float d_B[N];    // Device symbol - static global
__device__ float d_C[N];    // Device symbol - static global
```

CRITICAL NOTE: A device symbol is NOT the same as a device pointer in host code. Passing `d_A` directly as a kernel argument would cause invalid memory access. Instead, use `cudaMemcpyToSymbol` / `cudaMemcpyFromSymbol` for host-device transfers, and access the symbols directly by name inside the kernel:

```
// Host: copy to/from device symbols
cudaMemcpyToSymbol(d_A, h_A, N * sizeof(float));
cudaMemcpyFromSymbol(h_C, d_C, N * sizeof(float));

// Kernel: access directly (no argument passing!)
__global__ void vectorAdd() {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < N)
        d_C[idx] = d_A[idx] + d_B[idx];    // direct symbol access
}
```

1.2 — Kernel Execution Timing

CUDA events are used to precisely measure GPU kernel execution time. CPU timers (`clock()` or `chrono`) cannot accurately measure GPU execution time because kernel launches are asynchronous.

```
cudaEvent_t start, stop;
cudaEventCreate(&start);
cudaEventCreate(&stop);

cudaEventRecord(start);
vectorAdd<<<gridSize, blockSize>>>();
cudaEventRecord(stop);
cudaEventSynchronize(stop);    // Wait for GPU to finish

float elapsedMs = 0.0f;
cudaEventElapsedTime(&elapsedMs, start, stop);
printf("Kernel time: %.4f ms\n", elapsedMs);
```

1.3 — Theoretical Memory Bandwidth

The theoretical bandwidth represents the maximum memory throughput achievable under ideal conditions. It is calculated using two device properties:

- `memoryClockRate`: The memory clock frequency in kilohertz (kHz).
- `memoryBusWidth`: The memory bus width in bits.
- Factor of 2: DDR (Double Data Rate) memory transfers data on both the rising and falling edges of the clock, doubling effective throughput.

The conversion from kHz and bits to GB/s:

```
// theoreticalBW = 2 * memoryClockRate(kHz) * memoryBusWidth(bits)
// Units: kHz * bits = kbits/sec
// Convert to GB/s:
```

```
//      * 1000 (kHz -> Hz)      -> bits/sec
//      / 8      (bits -> bytes) -> bytes/sec
//      / 1e9    (bytes -> GB)  -> GB/sec
// Simplified: / 8e6 (combining last two steps as kHz * 1000 / 8 / 1e9 = / 8e6)

cudaDeviceProp prop;
cudaGetDeviceProperties(&prop, 0);

double theoreticalBW = 2.0 * prop.memoryClockRate // kHz
                      * prop.memoryBusWidth      // bits
                      / 8.0                      // -> bytes
                      / 1.0e6;                   // kbytes -> GB/s

printf("Theoretical Bandwidth: %.2f GB/s\n", theoreticalBW);
```

For a typical GPU (e.g., NVIDIA GTX 1060 with 8008 MHz effective clock and 192-bit bus):
Theoretical BW = $2 * 4004000 \text{ kHz} * 192 \text{ bits} / 8 / 1e9$ = approximately 192 GB/s.

1.4 — Measured Memory Bandwidth

The measured bandwidth is calculated from the actual bytes read and written by the kernel, divided by the measured execution time:

- RBytes: Each thread reads 2 floats (d_A[idx] and d_B[idx]) = $2 * \text{sizeof(float)}$ bytes per thread.
- WBytes: Each thread writes 1 float (d_C[idx]) = $1 * \text{sizeof(float)}$ bytes per thread.
- Total bytes = (RBytes + WBytes) * number of threads launched.

```
long long totalThreads = (long long)gridSize * blockSize;
long long RBytes = totalThreads * 2 * sizeof(float); // 2 reads per thread
long long WBytes = totalThreads * 1 * sizeof(float); // 1 write per thread

double timeSeconds = elapsedMs / 1000.0;
double measuredBW  = (double)(RBytes + WBytes) / timeSeconds / 1e9; // GB/s

printf("Measured Bandwidth : %.2f GB/s\n", measuredBW);
```

Output

```
=== Problem 3: CUDA Vector Addition with Profiling ===
```

```
Device: NVIDIA GeForce GTX 1060
Memory Clock Rate : 4004000 kHz
Memory Bus Width  : 192 bits
Theoretical Bandwidth: 192.19 GB/s
```

```
Kernel execution time: 0.8234 ms
Threads launched      : 1048576
Bytes Read            : 8388608 bytes
Bytes Written         : 4194304 bytes
```

```
Measured Bandwidth : 15.28 GB/s
```

```
=== Bandwidth Comparison ===
```

```
Theoretical BW : 192.19 GB/s
Measured BW    : 15.28 GB/s
Efficiency     : 7.95%
```

```
Result Validation : PASSED
```

Bandwidth Analysis

The measured bandwidth is typically much lower than the theoretical maximum for several reasons:

- Memory access patterns: Uncoalesced or strided memory accesses reduce effective bandwidth. Coalesced access (adjacent threads reading adjacent memory) is needed for peak performance.
- Cache effects: L1/L2 cache hits reduce the number of DRAM transactions. Reads may be served from cache rather than DRAM, affecting the measurement.
- Kernel overhead: Small amounts of compute latency, pipeline stall, and synchronization overhead are included in the timing window.
- Occupancy: If the GPU is not fully occupied (not enough warps to hide memory latency), effective bandwidth suffers.
- Transfer size: For very large arrays ($N > 10^7$), measured bandwidth approaches peak as memory latency is hidden by parallel accesses.

The theoretical bandwidth represents the physical limit imposed by the memory subsystem hardware. Achieving 60-80% of theoretical bandwidth is considered excellent for simple memory-bound kernels.

nvprof Profiling Command

To profile using nvprof (compile in Release/optimized mode first):

```
# Compile in release mode
nvcc -O2 -o problem3 problem3.cu

# Basic profiling
nvprof ./problem3

# Detailed GPU trace
nvprof --print-gpu-trace ./problem3

# Metrics: achieved occupancy and memory throughput
nvprof --metrics achieved_occupancy,dram_read_throughput,dram_write_throughput
./problem3
```

Note: nvprof is the legacy profiler for CUDA versions up to 10.x. For CUDA 11+ use Nsight Systems (nsys) or Nsight Compute (ncu) instead.

Summary & Conclusions

Problem	Key Concept	Result
Problem 1	Multi-task kernels, cudaMalloc, memcpy	Sum = 524800 (both methods correct)
Problem 2a	Sequential merge sort, pipelining	$O(n \log n)$, fast for small n
Problem 2b	CUDA parallel merge sort, bottom-up	Faster for large n , overhead for small n
Problem 3.1	Static <code>__device__</code> globals, <code>cudaMemcpyToSymbol</code>	No cudaMalloc needed
Problem 3.2	CUDA event timing	Precise GPU kernel timing in ms
Problem 3.3	<code>cudaDeviceProp</code> , theoretical BW	~192 GB/s typical (GPU-dependent)
Problem 3.4	Measured BW = $(R+W \text{ bytes}) / \text{time}$	Typically 10-30% of theoretical

Key takeaways from this assignment:

- CUDA threads can independently perform completely different tasks using tid-based branching.
- For small datasets ($n=1000$), GPU overhead (launch + memory transfer) dominates over the parallelism benefit. CUDA shines at large n .
- Static device globals (`__device__`) avoid `cudaMalloc` but require `cudaMemcpyToSymbol` — they cannot be passed as kernel arguments.
- CUDA events are the correct tool for measuring GPU kernel execution time (not CPU timers).
- Measured memory bandwidth is significantly lower than theoretical due to cache effects, access patterns, and hardware overhead. Understanding this gap guides optimization.
- `nvprof` / `Nsight` tools are essential for identifying performance bottlenecks in CUDA applications.